



Methods of Artificial Intelligence

1. Introduction to Local Search

Kai-Uwe Kühnberger, Nohayr Muhammad
Winter Term 2022/2023
November 4th, 2022

Today


- Solving problems by searching
- Search problems
- (Classical) Search algorithms
- Searching in complex environments
- Local search main ideas

Solving Problems by Searching

A quick recap!

Problem-Solving agent

An agent that can look ahead to find a sequence of actions that will eventually achieve its goal.



A four-phase problem-solving process:

- **Goal formulation**

Goals organize behaviour by limiting the objectives and hence the actions to be considered.

- **Problem formulation**

A description of the states and actions necessary to reach the goal—an abstract model of the relevant part of the world.

- **Search**

The simulation of sequences of actions in the model, searching until a sequence of actions that reaches the goal is found, if any.

- **Execution**

The agent can now execute the actions in the solution, one at a time.

Search problems examples

Route-planning

- Goal formulation
*Go from starting point A to point B.
Possibly with (shortest, fastest, gas saving,
...) path.*
- Problem formulation
 - *States: Locations (e.g. cities), time, ...*
 - *Actions: Drive left, right, ...*
 - *Possibly: A cost assigned to actions
(time, distance, fuel consumption, ...)*



Solving puzzles

■ Goal formulation


Starting from certain, possibly random, configuration to reach the desired configuration.

Possibly with least number of moves.


■ Problem formulation

- *States: Puzzle configuration*
- *Actions: Move tile, add number,...*
- *Possibly: A cost assigned to actions (time, number of moves, ...)*

5	3			7			
6			1	9	5		
	9	8				6	
8				6			3
4			8		3		1
7				2			6
	6					2	8
			4	1	9		5
				8		7	9



8		6
5	4	7
2	3	1



Robot motion planning

- Goal formulation
*Perform desired task, e.g. pick up something.
Possibly in the (fastest, safest, energy saving, ...) manner.*
- Problem formulation
 - *States: Position of robot parts*
 - *Actions: Translate, rotate parts*
 - *Possibly: A cost assigned to actions
(time, risk factor, energy consumption, ...)*



Definition: Search Problem

- A search problem is defined as a 4-tuple $\langle S, R, s_0, G \rangle$ with
- The search space (state space) S
 - A successor relation (transition relation) $R \subset S \times S$
 - An initial state $s_0 \in S$
 - A goal predicate $G : S \rightarrow \{true, false\}$ which terminates the search

Also, a possible transition cost.

Definition: Solution

A solution to such a search problem is a sequence of states

$$(s_0, \dots, s_n),$$

starting with initial state

$$s_0$$

with adjacent states in the transition relation

$$(s_i, s_{i+1}) \in R \text{ for } i = 0, \dots, n - 1$$

and ending at a goal state or

$$G(s_n) = \text{true}$$


A general search "algorithm"

A search algorithm takes a search problem as input and returns a solution, or an indication of failure.

This can be done by creating a search tree over the state-space graph, forming various paths from the initial state, trying to find a path that reaches a goal state.


State-Space Graph

The state space can be represented as a graph in which the vertices are states and the directed edges between them are actions.




Search Tree

- Each node in the search tree corresponds to a state in the state space.
- The edges in the search tree correspond to actions.
- The root of the tree corresponds to the initial state of the problem.




State Graph vs. Search Tree

State space graph



Search tree



A rough sketch of a general search algorithm:

- Generate a search tree:
 - Start with s_0 as root.
 - Add successors to a leaf-state s_i for all transitions $(s_i, s_j) \in R$
- A path from the root s_0 to a node s in the tree with $G(s) = \text{true}$ is a solution
- Traverse the search tree somehow to collect the solutions

There are different ways to traverse somehow → Different search algorithms.

Classical Search

☐ Uninformed (blind)

■ Random search

■ Systematic search

- ☐ Depth-first search (DFS)
- ☐ Breadth-first search (BFS)
- ☐ Variations
 - Uniform-cost search
 - Depth-limited search

■ Features

- ☐ no information about the path cost from the current state to the goal state

☐ Informed (heuristic)

■ Systematic search

- ☐ Greedy best-first search
- ☐ A* search

■ Features

- ☐ heuristic information (estimate) about the path cost from current state to goal state




From Classical Search to Local Search

Classical Search

- find (optimal) sequence of actions Sequence might be irrelevant
- that leads from initial state Might not exist
- to a goal state Might be hard to test

Local Search

- find „nearly“ feasible/optimal state



Local Search Applications

Local Search is applied if search spaces are too big for exhaustive search.

Typical problem types are:


- Constraint Satisfaction Problems
- (Discrete) Optimization Problems
 - Production planning (maximize profit)
 - Scheduling (minimize conflicts)
 - Machine Learning (minimize classification error)

States and Neighborhoods

- States often correspond to **variable assignments**

X1	X2	X3	X4
0	1	1	0

- **neighborhood** determines which assignments can be reached from given assignment



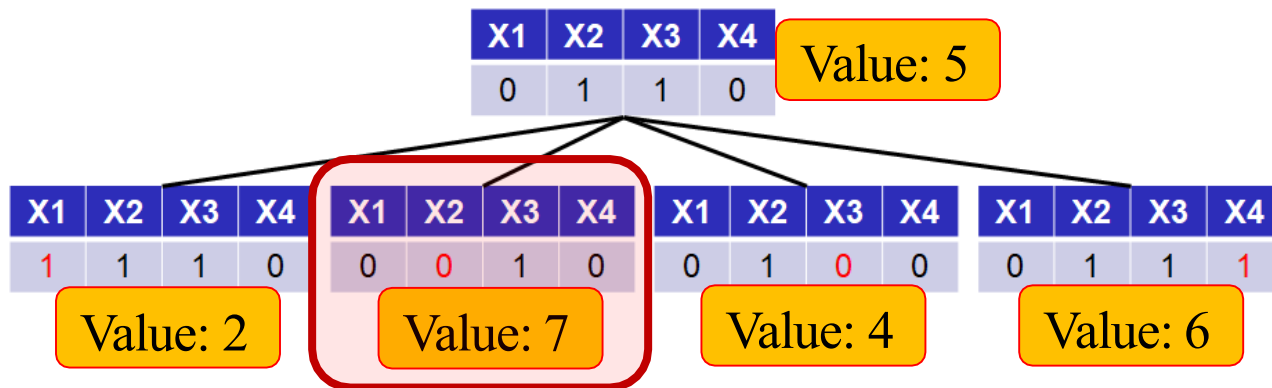
Local Search Main Ideas

Rough Idea

- start from some **initial assignment** (e.g. random choice)

X1	X2	X3	X4
0	1	1	0

- select a "promising" **neighbor** of the current assignment



- continue selecting neighbors until we cannot improve anymore


Solving local search problems

1. Define **state space**
2. Define **neighborhood**
3. Define **objective function**
(minimize $f(x)$ by maximizing $-f(x)$)
4. Apply **local search algorithm**

Local Search Example

How can we apply local search to typical problems?

N Queens Problem



1. **State-Space:** board configurations with one queen per column (tuple (8,3,7,4,2,5,1,6) corresponds to board configuration above)
2. **Neighborhood:** configurations that can be obtained by moving a single queen to another field in the same column

3. **Objective Function:** number of pairs of queens that can attack each other directly or indirectly (*maximize negative objective function*)

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18

- Queen in column 1 **can directly attack** queen in column 2
- Queen in column 1 **can indirectly attack** queen in column 3
- State has **value** 17
- Numbers show **values of neighbors**

4. Apply a local search algorithm

Next time...

We will discuss the following local search algorithms:

- Hill-Climbing
- Simulated Annealing
- Local Beam Search
- Genetic Algorithms

Summary

- Problem-solving by searching
- Formulating search problems
- Difference between classical and local search algorithms
- If we are only interested in a solution and not in the solution path,
Local Search can be better suited than Classical Search
- this is typically the case for **optimization** and **constraint satisfaction** problems

Slides are, mainly, Dr. Tobias Thelen slides.
Some presented ideas are adapted from other courses like CS188 at UC Berkeley.

Further Readings

Lecture is mainly based on:

*Russell, S., Norvig, P. Artificial Intelligence - A modern approach.
Pearson Education: 2010.*

More details on presented (and similar algorithms) can be found in:

*Burke, E. K., & Kendall, G. Search methodologies - Introductory Tutorials in
Optimization and Decision Support Techniques. Springer US: 2014.*