



Methods of Artificial Intelligence

5. Planning II: Probabilistic Planning

Nohayr Muhammad

Winter Term 2022/2023




November 30th, 2022

Last time..

- Classical Planning
- Find sequence of actions that transforms given state into goal state
- An action “guarantees” reaching certain states
- Situation Calculus
- STRIPS
- PDDL

Example

S1:



- `onTable (C, S1)`
- `on (B,C, S1)`
- `on (A,B, S1)`
- `clear (A, S1)`

← `put_on_table(A,S1)`

- `onTable (C, put_on_table (A, S1))`
- `on (B,C, put_on_table (A, S1))`
- `onTable (A, put_on_table (A, S1))`
- `clear (A, put_on_table (A, S1))`
- `clear (B, put_on_table (A, S1))`


← `put(B,A, put_on_table(A,S1))`

- `onTable (C, put (B,A, put_on_table (A, S1)))`
- `on (B,A, put (B,A, put_on_table (A, S1)))`
- `onTable (A, put (B,A, put_on_table (A, S1)))`
- `clear (B, put (B,A, put_on_table (A, S1)))`
- `clear (C, put (B,A, put_on_table (A, S1)))`

Example: State-Space Model

Methods of Artificial Intelligence WS 2022/2023






- Operator application spans a search space
- s_i is connected to s_j iff s_j results from s_i when applying a single operator



Today..

- Probabilistic Planning
- Dealing with uncertainty
- Dealing with long-term consequences
- Markov Decision Processes (MDPs)



Introduction

0:14h


Today Tue, 29.11.2022

17:33
+27






17:47
+27

Departure Bus M1: 17:34+27
Osnabrück Roopstraße
0:14h 0x
from 2.70 €


 Journey suggestion according to current traffic.

17:43
+19






17:57
+19

Departure Bus M1: 17:44+19
Osnabrück Roopstraße
0:14h 0x
from 2.70 €




 Journey suggestion according to current traffic.

18:03

18:17
+10

Departure Bus M1: 18:04+10
Osnabrück Roopstraße
0:14h 0x
from 2.70 €

Aktuelle Verkehrsmeldungen

Region:

Baden-Württemberg

Bayern

Berlin

Brandenburg

Bremen

Hamburg

Hessen

Mecklenburg-Vorpommern

Niedersachsen

Nordrhein-Westfalen

Rheinland-Pfalz

Saarland

Sachsen

Sachsen-Anhalt

Schleswig-Holstein

Thüringen

Verkehrsmittel:

Regionalverkehr

Fernverkehr

Streckensperrung:
Fernverkehr ist zwischen Berlin und Hannover bis voraussichtlich 16.12.2022 beeinträchtigt
Streckensperrung - bis vsl. Freitag, 16.12.2022

- In many situations, we have to deal with **uncertainty**:
 - Taking a particular road is often fast, but sometimes there is a traffic jam
 - An order is usually delivered within two days, but can be delayed due to labor strike or logistical mistakes
 - When navigating a robot, sensor information is inherently uncertain

- Furthermore, actions often have long-term consequences
 - Not recharging the battery saves time now
 - but we may run out of energy later
 - Canceling insurance increases our budget now
 - but we may lose a lot of money later
 - Extending maintenance intervals may decrease spendings now
 - but may result in business interruptions due to technical failures later

- Markov Decision Processes (MDP)
 - Basics
 - The Grid World Example
 - A Sample Run
 - Policies and Rewards
 - Computing Discounted Rewards
 - Playing with the Rewards
- Policy Iteration Algorithm
 - Policy Evaluation
 - An Example of Policy Evaluation
 - Policy Improvement

Markov Decision Processes

Basics

“Markov” generally means that given the present state, the future and the past are independent

For Markov decision processes, “Markov” means action outcomes depend only on the current state

$$\begin{aligned} P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0) \\ = \\ P(S_{t+1} = s' | S_t = s_t, A_t = a_t) \end{aligned}$$



Andrey Markov
(1856-1922)

This is just like search, where the successor function could only depend on the current state (not the history)

- **Markov Decision Processes (MDPs)** take account of both mentioned problems:
 - uncertainty
 - long-term consequences
- Intuitively, MDPs describe
 - **probability of state changes** caused by actions
 - **short-term rewards** of actions in particular states
- A **policy** determines in what state we choose what action
- We evaluate policies by their **expected reward** with respect to
 - uncertainty of outcome of actions and
 - future rewards



- a **Markov Decision Process** is a tuple (S, A, p, r) , where
 - S is a set of **states**
 - A is a set of **actions**
 - p is the **state transition probability function**

$p(s' | s, a)$: probability that performing action a in state s yields state s'

- r is the **reward function**



$r(s, a)$: short-term benefit of performing action a in state s

Example: Grid World



- Move agent from Start to one of the goal fields
- Agent can move in four directions with uncertain outcome
- Game ends when agent performs an arbitrary move to one of the goal fields (special terminal states)

Example: States and Actions




- **States** encode agent's current position

$$S = \{ (1,1), \dots, (4,1), (1,2), \dots, (4,2), (1,3), \dots, (4,3) \}$$

- **Actions** encode agent's possible actions

$$A = \{ \text{up, down, left, right} \}$$





- 20% chance for a transmission error that will drive robot in wrong direction
- **Transition function p :** when performing **action a** , with probability
 - 0.8, we will move in the intended direction
 - 0.1, we will move in direction 90° clockwise to intended direction
 - 0.1, we will move in direction 90° counterclockwise to intended direction

Unless action leads out of grid world or to an obstacle




– in this case, nothing happens (robot is blocked by wall)

Example: Transition Probabilities




- p : when performing **left**, with probability
 - 0.8, we will move left
 - 0.1, we will move up
 - 0.1, we will move down

Unless action leads out of grid world – in this case, nothing happens




- $p((1,2) \mid (1,1), \text{up}) = 0.8$
- $p((1,1) \mid (1,1), \text{left}) = 0.8 \text{ (left)} + 0.1 \text{ (down)} = 0.9$
- $p((1,3) \mid (1,1), \text{up}) = 0$
- $p((2,1) \mid (1,1), \text{up}) = 0.1$



- **Rewards** give incentive to reach desired goal state
- $r(s, a) = -0.04$ for all actions a and all states $s \neq (4,2)$, $s \neq (4,3)$
(each action that does not finish game, results in penalty)
- $r((4,2), a) = -1$ for all actions a
- $r((4,3), a) = 1$ for all actions a

Grid World Example: A sample run


Example: Sample Run



Cumulative Reward: 0

Example: Sample Run


move up (works as intended)



Cumulative Reward: -0.04

Example: Sample Run


move up (robot moves right instead)



Cumulative Reward: -0.08

Example: Sample Run


move up (works as intended)



Cumulative Reward: -0.12

Example: Sample Run


move right (works as intended)



Cumulative Reward: -0.16

Example: Sample Run


move right (works as intended)



Cumulative Reward: -0.2

Example: Sample Run


move right (robot moves down instead)



Cumulative Reward: -0.24

Example: Sample Run


move up (works as intended)



Cumulative Reward: -0.28

Example: Sample Run


move right (works as intended)



Cumulative Reward: -0.32

Example: Sample Run

move right (action does not matter – game ends)




Cumulative Reward: **0.68**

Markov Decision Processes

Policies and Rewards

- a **(deterministic Markov) Policy** assigns an action to each state
- Formally, a policy is a mapping π from S to A



- $\pi((1,1)) = \text{up}$
- $\pi((1,2)) = \text{up}$
- ...

- When following policy π , we can compute the probability $P(S_k = s)$ of being in state s in the k -th step

- If we have states s_1, \dots, s_n , the **expected reward** is

$$E_{\pi}[r_k] = P(S_k = s_1) * r(s_1, \pi(s_1)) + \dots + P(S_k = s_n) * r(s_n, \pi(s_n))$$

- The **expected reward** of policy π is given by the series

$$E_{\pi}[r_1] + E_{\pi}[r_2] + E_{\pi}[r_3] + E_{\pi}[r_4] + \dots \quad (\text{possibly infinite series})$$

- The expected reward of policy π is given by the series

$$E_{\pi}[r_1] + E_{\pi}[r_2] + E_{\pi}[r_3] + E_{\pi}[r_4] + \dots \quad (\text{possibly infinite series})$$

- It may be reasonable to **discount future rewards**
 - because future rewards may be less valuable (economics) or
 - to guarantee convergence of the series

- The **γ -discounted reward of policy π** is given by the series

$$E_{\pi}[r_1] + \gamma * E_{\pi}[r_2] + \gamma^2 * E_{\pi}[r_3] + \gamma^3 * E_{\pi}[r_4] + \dots \quad \text{where } \gamma \text{ is a real number}$$

between 0 and 1

- The γ -discounted reward of policy π is given by the series

$$E_{\pi}[r_1 + \gamma * r_2 + \gamma^2 * r_3 + \gamma^3 r_4 + \dots]$$

where γ is a real number between 0 and 1

- Which are the first four terms for $\gamma=0.5$?

$$r_1, 0.5 * r_2, 0.25 * r_3, 0.125 * r_4$$



(weight of rewards halves with every step in the future)

Playing with the Rewards

- We can compute an optimal policy with respect to the γ -discounted rewards using different approaches
 - Policy Iteration
 - Value Iteration
 - Linear Programming
- Worst-case runtime of all approaches is polynomial in the number of states and actions

Example: Grid World

- The following picture shows optimal policy and expected values ($\gamma=1$) when starting from different states



3	→	→	→	+1
2	↑		↑	-1
1	↑	←	←	←
	1	2	3	4

3	0.812	0.868	0.918	+1
2	0.762		0.660	-1
1	0.705	0.655	0.611	0.388
	1	2	3	4

- Note that to get **from start to positive goal**, we need at least 5 steps, no matter whether we go up or right in the first step

3	→	→	→	+1
2	↑		↑	-1
1	↑	←	←	←
	1	2	3	4


3	0.812	0.868	0.918	+1
2	0.762		0.660	-1
1	0.705	0.655	0.611	0.388
	1	2	3	4

- **Why does (1,2) have a higher value than (2,1)?**

When going right, there is a higher risk of ending up in negative goal due to uncertainty in moves

- **Why does optimal policy recommend going back to start from (2,1)?**

Since the expected value from start is 0.094 higher than from (3,1) it is worth going back.




3	0.812	0.868	0.918	+1
2	0.762		0.660	-1
1	0.705	0.655	0.611	0.388
	1	2	3	4


- **Why does optimal policy recommend going left from (3,1)?**

Because one should consider also unintended effects of actions: choosing *up* might result in a move to the right, and (4,1) has a bad expected value.

We'll see this more formally in a moment.



- The picture above shows optimal policy when steps from non-goal states are rewarded with -2 (penalty)




- **Why is it more reasonable to move right from the start state now?**

Each step is significantly more expensive than the reward that we can get from any goal state.

Therefore, the agent should try to end the game as soon as possible.

The negative goal state can be reached before the positive one.



- Reward -0.4

Policy Iteration Algorithm

- Our **main problem** is finding an optimal policy (plan)
- We can consider **subproblems of increasing difficulty**
 - Given MDP and policy, **evaluate policy**
 - Given MDP, **find optimal policy**
 - Find optimal policy for unknown MDP (**reinforcement learning**)

- Initialize policies $\pi(s)$ to random actions
- Repeat
 - **Step 1: Policy evaluation:** calculate value $V_{\pi}(s)$ for each s
 - **Step 2: Policy improvement:** update policy using one-step look-ahead
- Until policy doesn't change

Policy Evaluation

- How can we **evaluate deterministic policy** π ?
- Consider **value function** $v: S \rightarrow \mathbb{R}$ that maps states to values
- We let $v_\pi(s)$ be the **expected discounted reward** when starting in s and following policy π

$$v_\pi(s) = E_\pi[r_1 + \gamma * r_2 + \gamma^2 * r_3 + \gamma^3 * r_4 + \dots \mid S_0 = s]$$

- We **represent v by an array**
 $[v(s_1), v(s_2), v(s_3), \dots, v(s_n)]$
- We **initialize** all values with 0
 $[0, 0, 0, \dots, 0]$
- We then **update values** by adding up the reward for the next action (given by policy) and the current expected value of the next state
- One can show that this approach converges to the true values of π
- However, for $\gamma=1$, values may be unbounded (termination problem)

- **Input:** - MDP (S, A, p, r)
 - deterministic Policy π
 - discount factor γ

Output: - value function v_π

Initialize $v(s) = 0$ for all s in S

do

for each s in S




$$v(s) \leftarrow r(s, \pi(s)) + \gamma * \sum_{s' \text{ in } S} p(s' | s, \pi(s)) * v(s')$$

until change in v 'is negligible'


return v

Iterative Policy Evaluation Example

($\gamma=1$)






Initialization



Iterative Policy Evaluation Example




($\gamma=1$)




After first iteration

3	-0.04	-0.04	-0.04	1
2	-0.04		-0.04	-1
1	-0.04	-0.04	-0.04	-0.04
	1	2	3	4

$(\gamma=1)$






Second iteration




$$\begin{aligned}
 & -0.04 \\
 & + 0.8 * -0.04 \\
 & + 0.1 * -0.04 \\
 & + 0.1 * -0.04 \\
 & \hline
 & = -0.08
 \end{aligned}$$

Iterative Policy Evaluation Example

($\gamma=1$)






Second iteration




$$\begin{aligned}
 & -0.04 \\
 & + 0.8 * 1 \\
 & + 0.1 * -0.04 \\
 & + 0.1 * -0.04 \\
 & \hline
 & = 0.752
 \end{aligned}$$

$(\gamma=1)$



Second iteration




$$\begin{aligned}
 & -0.04 \\
 & + 0.8 * -0.04 \\
 & + 0.1 * -0.04 \\
 & + 0.1 * -1 \\
 & \hline
 & = -0.176
 \end{aligned}$$

Iterative Policy Evaluation Example

$(\gamma=1)$


- Here are again the final values of the (optimal) policy



3	0.812	0.868	0.918	+ 1
2	0.762		0.660	- 1
1	0.705	0.655	0.611	0.388
	1	2	3	4

$(\gamma=1)$

- Value function is stable under updates



$$\begin{aligned}
 & -0.04 \\
 & + 0.8 * 1 \\
 & + 0.1 * 0.918 \\
 & + 0.1 * 0.66 \\
 & \text{-----} \\
 & = 0.9178
 \end{aligned}$$

Policy Improvement

- Now we know how to evaluate a given deterministic policy
- How can we **find an optimal policy**?
- Roughly speaking, we can do so by
 - Starting from an arbitrary policy
 - **Evaluating** the policy
 - **Improving** the policyAnd **iterating** until no improvement is possible anymore
- This approach is called **Policy Iteration**

How to improve Policies

- How can we **improve our policy**?
- One can show that an optimal policy π must be **greedy**

$$\begin{aligned} & r(s, \pi(s)) + \gamma * \sum_{s' \in S} p(s' | s, \pi(s)) * v(s') \\ &= \max_{a \in A} r(s, a) + \gamma * \sum_{s' \in S} p(s' | s, a) * v(s') \end{aligned}$$

for all states s

- If $\pi(s)$ is not greedy, we can improve policy by replacing $\pi(s)$ with a greedy action (**Policy Improvement Theorem**)
- If π is greedy, then it is optimal (**Bellman Optimality Equation**)

- **Input:** - MDP (S, A, p, r)
 - - discount factor γ
- **Output:** - optimal policy π

Initialize value estimate v and policy π arbitrarily

do

$v \leftarrow \text{evaluate } \pi$

(perform policy evaluation)

for each s in S

$\pi(s) \leftarrow \text{select greedy action with respect to } v$

until policy is stable (does not change anymore)

return π

Initialize value estimate v and policy π arbitrarily

do

$v \leftarrow \text{evaluate } \pi$

(perform policy evaluation)

for each s in S

$\pi(s) \leftarrow \text{action } a \text{ that maximizes}$

$$r(s,a) + \gamma * \sum_{s' \text{ in } S} p(s' | s, a) * v(s')$$

until policy is stable (does not change anymore)

return π

- **Policy iteration is guaranteed to converge** to an optimal policy under mild assumptions
- However, there are again **some subtleties**
 - for $\gamma=1$, values may be unbounded (evaluation may not terminate)
 - algorithm may cycle between actions that yield equal values
 - (in particular, optimal policy may not be unique)
- In **Modified Policy Iteration**, we perform only a fixed number of evaluation steps before improving policy
 - may converge faster
 - can avoid divergence problems for $\gamma=1$

Summary

- **Markov Decision Processes** take account of
 - Uncertainty and
 - Long-term consequences
- An MDP (S, A, p, r) consists of a definition of states, actions, transition probabilities and rewards
- **policies** describe in what state we choose what action
- policies can be evaluated by their **expected reward**
- **planning** comes down to computing an optimal policy

The presented slides are mainly Dr. Marco Volpe's slides

Most topics can be found in:

Russell, S., Norvig, P. *Artificial Intelligence - A modern approach*.
Pearson Education: 2010.
(3rd Edition. Sections 17.1-17.3)

More **detailed information** can be found in:

LaValle, S. M. *Planning algorithms*.
Cambridge university press: 2006

Thrun, S., Burgard, W., & Fox, D. *Probabilistic robotics*.
MIT press: 2005.