



Support Vector Machines

Kai-Uwe Kühnberger
University of Osnabrück
15.01.2021

Machine Learning Methodology: Schedule

- Overview Machine Learning Sessions
 - Basics (last week)
 - Machine Learning
 - Important Concepts
 - Clustering Methods
 - Properties of Hypotheses
 - Classification Methods (today)
 - Support Vector Machines
 - Example: Document Classification
 - Classification Methods (next week)
 - Decision Trees
 - Random Forests
 - Literature

Idea

- Support Vector Machines (SVMs)
- Motivation
 - Single-layer neural networks have efficient learning algorithms but can only learn linear functions
 - Multi-layer neural networks can learn non-linear functions, but are hard to train
 - Idea: SVMs map an input space into a space of higher dimensionality:

Then separability is possible with linear functions

Support Vector Machines

If no simple linear hyperplane exists: map to a higher dimensional space

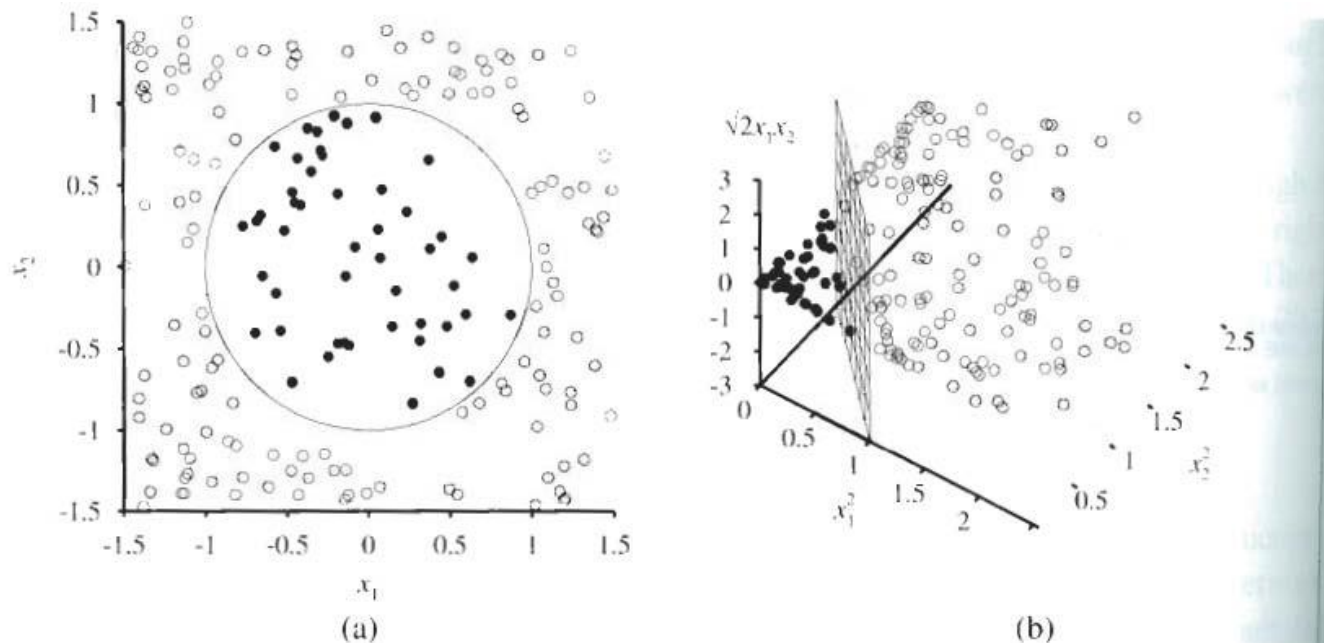



Figure 20.27 (a) A two-dimensional training with positive examples as black circles and negative examples as white circles. The true decision boundary, $x_1^2 + x_2^2 \leq 1$, is also shown. (b) The same data after mapping into a three-dimensional input space $(x_1^2, x_2^2, \sqrt{2}x_1x_2)$. The circular decision boundary in (a) becomes a linear decision boundary in three dimensions.

Russell &
Norvig, p. 747

Support Vector Machines

- Basic Ideas: optimal classifiers in case of linear separable data
- Simplest case: one-dimensional data



Support Vector Machines


- Remarks on support vector machines:
 - Supervised learning model for classification
 - Mathematical foundations were proposed by Vapnik and Chervonenkis in the early 1960ies.
 - SVM were introduced by Vapnik and Lerner 1963 (special case of support vector machine algorithm).
 - Later SVMs were extended in various directions, e.g. to cover also non-linear classifiers.
- SVMs are a de facto standard in machine learning.
 - They were applied to a broad variety of domains and showed often excellent results.
 - Examples: face detection, text categorization, bioinformatics, character recognition, prediction of financial time series, prediction of health risks etc.
 - For some of these applications deep learning approaches are now outperforming SVMs, e.g. face detection, character recognition
 - For others the situation is not as clear.

Basics

- Assume data is given by vectors of dimensionality n .
 - We presuppose there is an n -dimensional vector space given.
- Find a hyperplane that separates positive from negative examples
 - A hyperplane in an n -dimension space is an $n-1$ dimensional subspace.
- If this is possible, we can use a linear classifier.
- Because there are potentially many such hyperplanes, we choose the one that maximizes the distance between the hyperplane and the nearest positive and negative examples.

Basics

- Formally we have:
 - $(\mathbf{x}_1, \mathbf{y}_1), (\mathbf{x}_2, \mathbf{y}_2), \dots, (\mathbf{x}_n, \mathbf{y}_n)$ training data – class membership pairs
 - $y_i \in \{+1, -1\}$
 - Find a function $f: \mathcal{R}^n \rightarrow \{+1, -1\}$, such that $f(\mathbf{x}_i) = y_i$
 - Unseen data is then mapped via f .
- Assume again \mathcal{R}^n is given.
A separating hyperplane H is defined as follows:
$$H = \{\mathbf{x} \in \mathcal{R}^n \mid \langle \mathbf{w}, \mathbf{x} \rangle + b = 0\}$$
- Here: $\mathbf{w} \in \mathcal{R}^n$ orthogonal to H and $b \in \mathcal{R}$.



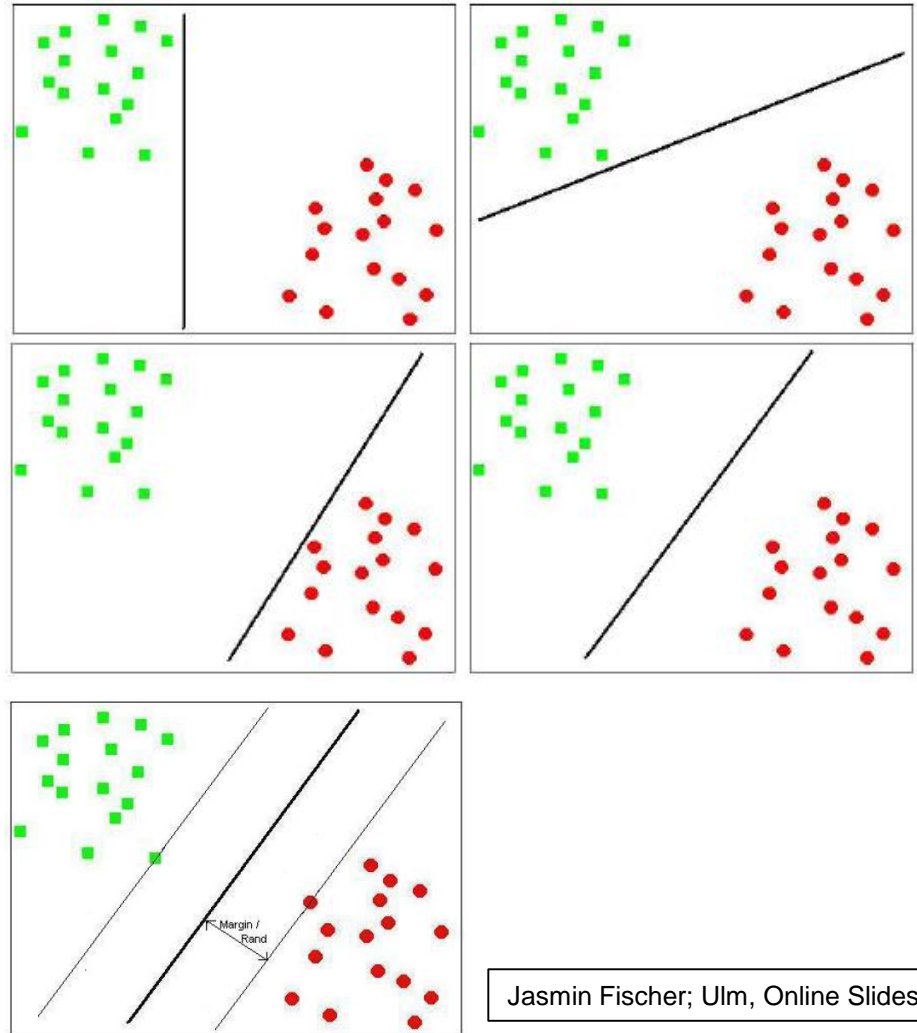
Jasmin Fischer; Ulm, Online Slides

Generalized Portrait

- The generalized portrait version by Vapnik & Lerner, 1963:
- A(n) (affine) hyperplane in H is defined by a weight vector $\mathbf{w} \in H$ and a bias $b \in \mathbb{R}$ with the condition that $\langle \mathbf{w}, \mathbf{x} \rangle + b = 0$.
 - All $\mathbf{x} \in X$ that satisfy the condition above are on the hyperplane.
 - If $\langle \mathbf{w}, \mathbf{x} \rangle + b < 0$, \mathbf{x} is a negative example, if $\langle \mathbf{w}, \mathbf{x} \rangle + b > 0$, \mathbf{x} is a positive example.
- The original generalized portrait algorithm computes the optimal hyperplane in H that separates positive and negative examples (if they are linearly separable).
 - Task: Maximize the distance between the hyperplane and the closest points (margin).

Generalized Portrait

- Assume training data is linearly separable
- What is an optimal hyperplane separating the positive and negative examples?
- Optimality means here maximizing the distance between the hyperplane and the closest points (margin).



Jasmin Fischer; Ulm, Online Slides

Generalized Portrait

- If the distance of two parallel hyperplanes that separate the classes is maximal, then the region between these hyperplanes is the margin.
- These hyperplanes can be specified as follows:


$$\langle w, x_1 \rangle + b = +1$$

$$\langle w, x_2 \rangle + b = -1$$

$$\Rightarrow \langle w, (x_1 - x_2) \rangle = 2$$

$$\Rightarrow \langle \frac{w}{\|w\|}, (x_1 - x_2) \rangle = \frac{2}{\|w\|}$$

- Margin is specified: $2/\|w\|$



Generalized Portrait

- Maximizing the margin is equivalent to minimizing the following expression.

$$\tau(\mathbf{w}) = \frac{1}{2} \cdot \|\mathbf{w}\|^2 \quad \text{with condition } y_i \cdot (\langle \mathbf{x}_i, \mathbf{w} \rangle + b) \geq 1$$

- If the objective function is found and the constraint holds that all training examples are classified correctly, finding the optimal hyperplane can be phrased as a constraint **optimization problem**.
- Such optimization problems can be solved by so-called Lagrange multipliers.
 - Formally: Consider the simplified optimization problem:

$$\begin{aligned} &\text{maximize } f(x, y) \\ &\text{subject to the condition } g(x, y) = 0. \end{aligned}$$

- We assume that both f and g have continuous first partial derivatives. We introduce a new variable λ called a **Lagrange multiplier** and study the **Lagrange function** (or **Lagrangian**) defined by

$$L(x, y, \lambda) = f(x, y) - \lambda \cdot g(x, y)$$

Generalized Portrait

- Transferring Lagrange multipliers and the Lagrange function to the case of support vector machines, we get the following:

$$L(w, b, \alpha) = \frac{1}{2} \cdot \|\mathbf{w}\|^2 - \sum_N \alpha_i \cdot (y_i \cdot (\langle \mathbf{x}_i, \mathbf{w} \rangle + b) - 1)$$

with Lagrange multipliers $\alpha = (\alpha_1, \alpha_2, \dots, \alpha_N)$


- Maximizing Lagrange multipliers α_i while minimizing \mathbf{w} and b , yields the following equation:
 - $\mathbf{w} = \sum_{i \in N} \alpha_i \cdot y_i \cdot \mathbf{x}_i$
 - Intuitively: the weight vector can be expressed as a linear combination (sum) of the training examples.

Generalized Portrait

- In fact, there is a rather strong consequence:
 - According to the Karush-Kuhn-Tucker saddle point condition:
 $\alpha_i \cdot (y_i \cdot (\langle \mathbf{x}_i, \mathbf{w} \rangle + b) - 1) = 0$
 - All points where $\alpha_i > 0$ specify the optimal hyperplane H (points closest to the hyperplane, i.e. support vectors).
 - The rest of the points ($\alpha_i = 0$) do not have an influence.

- In fact, the support vectors (points that are closest to the hyperplane) suffice to compute the weight vector as the linear combination above:

$$\mathbf{w} = \sum_{i \in N: \mathbf{x}_i \text{ Support Vector}} \alpha_i \cdot y_i \cdot \mathbf{x}_i$$



Support Vector Machines

- The new decision function can now be expressed as follows:

$$f(\mathbf{x}) = \text{sgn}(\sum_i \alpha_i \cdot y_i \cdot \langle \mathbf{x}, \mathbf{x}_i \rangle + b)$$

- Support vector machines allow to replace the dot product by a kernel (kernel trick):

$$f(\mathbf{x}) = \text{sgn}(\sum_i \alpha_i \cdot y_i \cdot k(\mathbf{x}, \mathbf{x}_i) + b)$$

- The kernel function $k: X \times X \rightarrow \mathbb{R}: (\mathbf{x}, \mathbf{x}') \rightarrow k(\mathbf{x}, \mathbf{x}')$ can be used to map a low-dimensional input space X to a feature space H with higher dimensionality:
 - An appropriate choice of the feature space H allows to linearly separate examples that are not linearly separable in the input space X .
 - A kernel allows to compute the decision boundary without computing the whole function but only the support vectors.

Kernel

- Again: The idea is to use a similarity measure (kernel)
$$k: X \times X \rightarrow \mathbb{R}: (x, x') \rightarrow k(x, x')$$
- Formal definition: A mapping $k: X \times X \rightarrow \mathbb{R}$ is called *kernel*, if there is a product space $(H, \langle \cdot, \cdot \rangle)$ and a mapping $\Phi: X \rightarrow H$, such that $k(x, x') = \langle \Phi(x), \Phi(x') \rangle$.
- Remarks:
 - X can be any set on which a kernel can be defined.
 - Training data is given by pairs (x_i, y_i) for $x_i \in X$ and $y_i \in \{0, 1\}$
 - An example for a kernel is the dot product $\langle \mathbf{a}, \mathbf{b} \rangle$ in a product space: $\langle \mathbf{a}, \mathbf{b} \rangle = a_1 b_1 + a_2 b_2 + \dots + a_n b_n$
 - Take, for example, a vector space for the dot product space, and vectors \mathbf{a} and \mathbf{b} as factors.

Kernels

- There are many possibilities to define kernels.
- The choice of the kernel is dependent on the learning task and the input domain.
- Here are three important kernels:

Kernel	Definition	Parameter
Linear Kernel	$k(x, x') = \langle x, x' \rangle$	none
Polynomial Kernel	$k(x, x') = \langle x, x' \rangle^d$	Polynomial degree d
Gaussian Kernel / RBF Kernel	$k(x, x') = \exp(-\frac{\ x-x'\ ^2}{2\sigma^2})$	Free parameter σ

Klaus-Michael Lux, Bachelor Thesis

- For linear separable problems, non-linear kernels do not increase performance.
- For various applications new kernels were invented.

Support Vector Machines

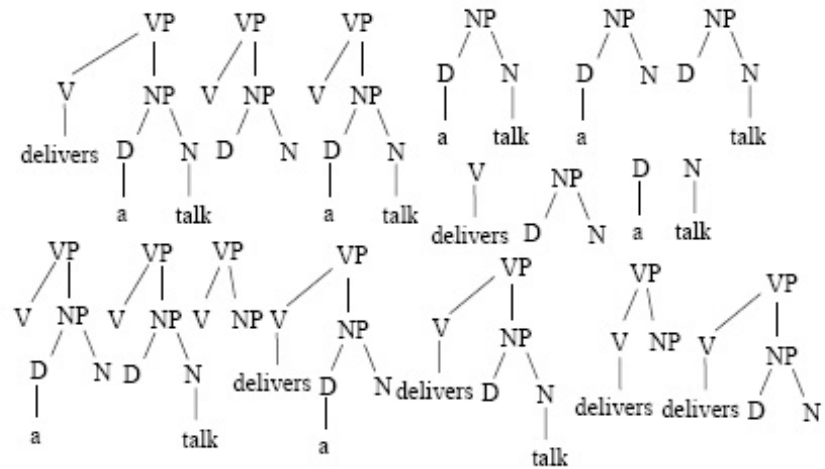
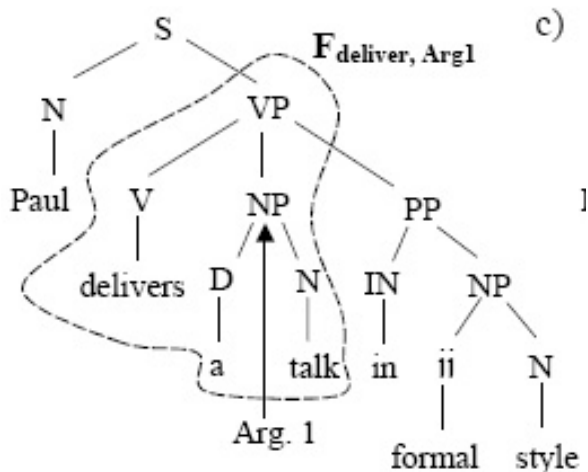
- Support vector machines are defined on binary classification problems.
 - How can multi-class classification problems be addressed?
- Here are two methods that have been proposed:
 - One-against-all
 - For k classes, k binary SVM models are trained. The i th SVM is trained using the examples in the i th class as positive examples and all other examples as negative examples.
 - For classifying x , feed x it into all k models. Select the model for which the decision function is maximal.
 - One-against-one
 - For k classes, $k \cdot (k-1)/2$ SVM models are trained. Each is trained to decide between two particular classes, leaving out all examples from other classes.
 - Each SVM votes for one class. The class with the most votes is selected.

Example

Document Classification

Example (based on Geibel et al., 2007)

- We want to classify DOM (document object model) trees of documents
 - DOM trees are rooted, labeled, ordered trees.
 - Inner nodes are labeled with XML tags, leaves might be labeled with sentences (maybe also represented as trees).
 - Approach works for every tree-like presentation.
- The Parse Tree Kernel (Collins & Duffy)
 - A tree is described by all possible intermediate parse trees t_i resulting in a feature vector $(\phi_{t_1}, \dots, \phi_{t_k})$.



Example

- The **Parse Tree Kernel (Collins & Duffy)**:
 - Assume $\Delta(v, v')$ is defined as the number of isomorphic mappings of partial parse trees rooted in nodes v and v' .
 - A kernel can be computed recursively by $k(T, T') = \sum_{v \in V, v' \in V} \Delta(v, v')$ such that
 - $\Delta(v, v') = 0$ if the productions applied in v and v' are different.
 - $\Delta(v, v') = 1$, if the productions in v and v' are identical and both nodes are pre-terminals.
 - For other non-terminals with identical productions (v_i is the i th child of v and $n(v)$ denotes the number of children of v):

$$\Delta(v, v') = \prod_{i=1}^{n(v)} (1 + \Delta(v_i, v'_i))$$

- This idea can now be applied to DOM trees.

Example

- The **Simple Tree Kernel SimTK** for DOM trees:
 - Incorporate kernel k^Σ operating on pairs of node labels (tags, attributes, text).
 - If there are either no children, or the number of children differs we set:

$$\Delta_{\text{SimTK}}(v, v') = \lambda \cdot k^\Sigma(\alpha(v), \alpha(v'))$$

- Else:
$$\Delta_{\text{SimTK}}(v, v') = \lambda \cdot k^\Sigma(\alpha(v), \alpha(v')) \left(1 + \prod_{i=1}^{n(v)} \Delta_{\text{SimTK}}(v_i, v'_i)\right)$$

α is a mapping from nodes to node labels Σ , λ is a parameter.

- Pros of the SimTK:
 - No grammar is presupposed.
 - We can include complex node labels, e.g. text in leave nodes.
- Shortcomings of the SimTK:
 - If the number of children differs, then the children are not compared.
 - If the number of children is not different, then they are only compared in the original order.

Example

- The **Left Aligned Tree Kernel** (compares just as many children as possible, if the number of children differ)
- Recursive case:

$$\Delta(v, v') = \lambda \cdot k^{\Sigma}(\alpha(v), \alpha(v')) \left(1 + \sum_{k=1}^{\min(n(v), n'(v'))} \prod_{i=1}^k \Delta(v_i, v'_i) \right)$$

- Shortcomings: Trees occurring more on the left have a higher influence than trees occurring on the right and no permutations are allowed.
- The Set Tree Kernel (treat children as a set)

$$\Delta(v, v') = \lambda \cdot k^{\Sigma}(\alpha(v), \alpha(v')) \left(1 + \sum_{i=1}^{n(v)} \sum_{i'=1}^{n'(v')} \Delta(v_i, v'_{i'}) \right)$$

- Shortcomings: no information about ordering retained at all.

Example

- The **Soft Tree Kernel**
- Idea: use a fuzzy / soft comparison of node positions using an RBF kernel:

$$k_{\gamma}(x, y) = e^{-\gamma(x-y)^2}$$

- Recursion:
$$\Delta(v, v') = \lambda \cdot k^{\Sigma}(\alpha(v), \alpha(v')) \cdot k_{\gamma}(\mu(v), \mu'(v')) \cdot \left(1 + \sum_{i=1}^{n(v)} \sum_{i'=1}^{n'(v')} \Delta(v_i, v'_{i'})\right)$$
- $\mu(v_i) = i$ specifies the position of child v_i of some node v , γ is a parameter, x and y are positions of children of a node.
- Pros: Has everything that is necessary.

Example

- Apply these kernels to artificial data and real data.
 - The class 1 examples all have a left-aligned subtree of the form $g(a, b(e), c)$.
 - The class 2 examples all have a general ordered subtree of the form $g(c, b, e(a))$, where gaps are allowed but the ordering of the subtrees c , b and $e(a)$ has to be preserved.
 - The class 3 examples contain subtrees of the form $g(c, b, a(e))$, where the child trees c , b and $a(e)$ are allowed to occur reordered and gaps might have been inserted, too.

Table: Optimal F-Measures

	Class 1	Class 2	Class 3
TagTK	0.727	0.6	0.736
LeftTK	0.909	0.363	0.44
SetTK	0.952	1.00	1.00
SoftTK	1.0	1.0	1.0
StringTK	1.0	1.0	1.0

Example

Class 1:

$f(n, h(m), g(a, b(e), c))$
 $f(h(m, g(a, b(e), c)))$
 $f(g(a, b(e), c, n), h(m))$
 $f(g(a, b(e), c, n, m))$
 $f(a, m, m, h(h(g(a, b(e), c))), n)$
 $f(g(a, b(e), c), g(e(a), c, a))$
 $f(h(m, c), g(a, b(e), c), g(m, n, m))$
 $f(g(a, b(e), c, h(h(m), n, b)))$
 $f(h(h(g(a, b(e), c))))$
 $f(g(a, b(e), c), a)$

Class 3:

$f(n, g(c, n, b, m, a(e)), h(m))$
 $f(m, e, b, h(g(b, n, c, m, a(e))))$
 $f(h(h(a(e), m, b, n, n, c)), b, e)$
 $f(e, b, g(m, c, a, b, e, e, a(e)))$
 $f(b, e, g(a(e), m, b, h(a), m, c))$

Class 2:

$f(n, h(n), g(c, m, b, n, e(a)))$
 $f(h(h(m), n, b), g(c, b, n, m, e(a)))$
 $f(h(g(h(n), c, b, e(a))), b)$
 $f(h(g(n, c, b, e(a))), h(m, b, h(a, n)))$
 $f(g(b(e), c, a), g(c, n, b, e(a)))$
 $f(g(c, h(c, n), b, h(h(h(b))), e(a)))$
 $f(g(c, b, e(a)), g(m, b, n, a, b(e), n, c))$
 $f(b, g(c, b, e(h(b, m, a), a)))$
 $f(g(g(a), c, d, m), g(c, b, h(n), e(a)))$
 $f(g(m, c, b, e(a), h(a)))$


Class 3 (continued):

$f(g(h(m, e), a(e), h(b), c, a, b))$
 $f(g(b, h(m), c, h(n), a(b, e, a)))$
 $f(g(m, h(n, h(n)), g(b, e, c, a(e))))$
 $f(g(a(e), g(h(n)), b, c), b, n)$
 $f(h(h(e), h(n), b, h(n), m), g(c, m, n, b, a(e)))$

Example

- Real world example:
- Approx. 35,000 short texts from Süddeutsche Zeitung were considered.
 - 31 Classes:
 - Bühnentip (theater)
 - Hochschulnachrichten (university news)
 - Fragen und Antworten (questions and answers)
 - Inhalt (content)
 - Wochenchronik (chronicle of the week)
 - Etc.
- Document DOM tree + class

Example



- Result for LeftTK (left aligned TK) are depicted.
- Many classes are learnable, some still have problems.
- SoftTK and SetTK: rather bad results.

Bibliography

- Support Vector Machines
 - Schölkopf, Smola: *Learning with Kernels*, MIT Press, 2001.
 - Vapnik and Lerner: Pattern Recognition Using Generalized Portrait Method. *Automation and Remote Control*, Vol. 24, 1963, pp. 774-780.
 - Vapnik, Vladimir N (2000). *The Nature of Statistical Learning Theory*. Springer.
 - Geibel, Gust, Kühnberger (2007): Variants of Tree Kernels for XML Documents. In: A. Gelbukh and A.F. Kuri Morales (Eds.): *MICA 2007*, LNAI 4827, pp. 850–860, 2007.